MR-NWM: Language-Model-Centered World Models for Multimodal, Multi-Rate RAN Control

Research preprint, October 2026

Abstract

MR-NWM is a language-model-centered architecture for multimodal, multi-rate radio access network (RAN) world modeling. A pretrained language transformer serves as the shared slow backbone. Continuous modality adapters receive state/traffic sequences, real/imaginary CSI and IQ features, topology nodes, and optional patch features, with explicit observation ages and sampling periods. Two causal views of the same language backbone separate operational preferences from exogenous physical evidence. A context-conditioned recurrent subnetwork then predicts numerical state transitions at the fast rate using a cached modulation packet, with optional slow/fast physics-expert sharing. We supply this pretrained-model integration alongside a lightweight, train-from-scratch reference variant. Three-seed, eight-variant measurements on a fully specified two-link queue/interference simulator evaluate only the lightweight variant: dynamics, distribution shift, controlled-language parsing, action effects, decision regret, and host-observed timing. Those results do not establish performance of a trained language-foundation RAN model. The project provides executable multimodal/gradient/causality tests and a reproducible path from the small reference experiment toward the intended language-based system, without claiming calibrated radio behavior or deterministic real-time deadlines.

1 Introduction

Network operators express objectives in terms of service quality, energy, reliability, and business priorities. A controller must translate these objectives into measurable quantities, evaluate candidate actions against network dynamics, and respect actuator and timing constraints. Language models can provide an extensible interface to these tasks, while continuous-valued predictors can represent the physical consequences of actions without converting every measurement into text.

Two broad designs are possible: a common multimodal backbone can process language and network measurements, or a semantic component can invoke a separate dynamics model. Neither design inherently determines the numerical representation of wireless observations. In particular, a language-derived backbone can receive continuous embeddings, and a tool interface can transmit continuous tensors. Nor does a modular interface necessarily prohibit joint training: differentiability depends on the operations and training graph. The relevant questions are therefore empirical: what should be shared, which information should influence the transition law, and how should computational work be divided across rates?

The full MR-NWM architecture places a pretrained language model at the center of slow multimodal fusion and intent processing. Its numeric adapters and fast physical subnetwork are jointly trainable around that backbone. The lightweight reference hypothesis is that a slow context encoder and a small recurrent transition operator can be coupled through weight sharing and low-dimensional conditioning, while keeping intent-to-objective compilation explicit. This construction builds on established multi-rate recurrence, hypernetworks, feature modulation, and low-rank parameterization [2, 3, 4, 5]. We do not claim these components individually as new. The contribution of this reference study is their concrete composition with an enforced preference/dynamics boundary and an evaluation protocol that can falsify benefits from sharing or gradient coupling.

The implementation makes three choices that are particularly consequential. First, continuous physical state and exogenous context enter through numeric encoders rather than text quantization. Second, the slow physical stream cannot read preference tokens, including through the recurrent expert. Third, a forecast spanning two macro windows passes its predicted terminal state and fast summary into the next window. Thus long-horizon errors cannot be hidden by resetting the initial condition. These choices are independently covered by regression tests.

2 Language-Model-Centered Multimodal Architecture

2.1 Foundation and modality-specific encoders

The full MR-NWM design uses a *pretrained language transformer as the shared slow backbone*. Continuous observations enter this backbone through trainable modality adapters rather than being converted into strings. A separate fast physical subnetwork performs the repeated numeric transitions between backbone updates. This separation preserves a language-based foundation while

assigning high-frequency prediction to a smaller computational path.

Let $x_T^{(m)}$ be observations of modality m, $v_T^{(m)}$ their validity mask, $\delta_T^{(m)}$ their observation ages, and ρ_m the acquisition period. A modality encoder creates continuous tokens

$$e_T^{(m)} = E_m(x_T^{(m)}) + e_{\text{type}}^{(m)} + E_{\text{time}}[\log(1+\delta_T^{(m)}), \log(1+\rho_m)]. \quad (1)$$

The implementation supports variable token counts, explicitly masked missing observations, and an input-token budget. A modality can be absent entirely. Negative ages, non-finite valid measurements, and malformed masks are rejected. These checks make the temporal information explicit; they are not a distributed clock-synchronization or missing-data recovery system.

For graph inputs, an optional normalized adjacency operator mixes neighboring node embeddings before fusion. For CSI/IQ, the caller supplies explicit real and imaginary components with known dimensions. Image or map patches can use the same registered-adapter interface, but this release does not include a pretrained vision encoder or raw-camera pipeline. Adapter presence alone does not establish useful cross-modal alignment. That alignment must be learned and tested on synchronized observation/action trajectories.

2.2 Two views of one language backbone

Physical predictions must be able to use information about the external world without confusing it with the operator's preferences. The same pretrained transformer F_θ therefore processes two separate input views. Let

$$Z_T^p = [E_s(s_{T,0}); E_c(c_T); E_H(H_{T-1}); \{e_T^{(m)}\}_m; E_K(K_T)], \quad (2)$$

where K_T contains externally supplied, validated evidence, such as an actual configuration change. A learned query q_p reads the physical view, while intent tokens are appended only to the semantic view:

$$h_T^p = P_p F_\theta([Z_T^p; q_p])_{\text{last}}, \quad (3)$$

$$h_T^I = P_I F_\theta([Z_T^p; E_I(I_T)])_{\text{last valid}}. \quad (4)$$

Both views use the *same language-model parameters*, with independent forward passes and masks. Preferences cannot enter the physical forward pass through attention or hidden-state reuse. By contrast, changes to observed modalities or exogenous evidence may legitimately change predicted dynamics. The caller is responsible for distinguishing a requested change from a change that has actually been applied or externally observed.

The projected physical context passes through the recurrent physics expert and then generates the fast-tier FiLM/low-rank packet. The physics expert can share its parameters with the fast recurrence. This is a coupling around the language trunk; the code does not claim to have converted the pretrained transformer's internal feed-forward layers into wireless SSM experts. The objective head uses h_T^I. Continuous state and structured objective heads are implemented; general natural-language explanations and arbitrary protocol action decoders remain additional output heads to develop and evaluate.

The `LanguageBasedWorldModel.from_pretrained` entry point loads an actual Hugging Face transformer checkpoint. A concrete small example is `Qwen/Qwen2.5-0.5B` [11]. Pretrained weights can be frozen while gradients still propagate *through* the transformer into numeric adapters. Disabling autograd around the frozen trunk would break that training path and is deliberately avoided. The numeric adapters, context projections, modulation heads, and dynamics predictor start untrained; pretrained text weights do not automatically provide a calibrated wireless world model.

2.3 Multiple observation rates and execution rates

The architecture distinguishes three quantities: each modality's acquisition period ρ_m, the slow backbone's refresh interval ΔT, and the fast predictor's step interval Δt, with $\Delta T = K\Delta t$ in the basic training graph. Observation tokens retain their own ages and periods even when combined at one slow refresh. A token budget requires explicit pooling or resampling before excessively long streams enter the language backbone.

Only two learned execution rates are implemented: one language-backbone update per macro window and K lightweight recurrent steps. Operator-level knowledge retrieval, long-horizon task decomposition, and model lifecycle management can run at a still slower external rate, but are not a third implemented neural controller in this release. The fast API does not call the language backbone; the contract tests count backbone calls to enforce this property.

Placement in Non-RT management, a Near-RT application, or an in-node predictor depends on measured end-to-end budgets. No general language-model size or Python module name establishes that placement. In particular, the language-foundation version's latency must be measured separately from the lightweight prototype's timing table. RF/CSI sampling, tensor fusion, model inference, network transport, and actuator execution contribute different parts of the overall budget.

2.4 What the two code paths establish

The `foundation.py` path implements the intended language-centered architecture and registered multimodal token adapters. Its tests cover causal preference separation, graph/missing-data handling, input ages, frozen-trunk gradient flow, and cached fast execution. The optional integration script records the checkpoint identity and resolved revision. We executed this check with

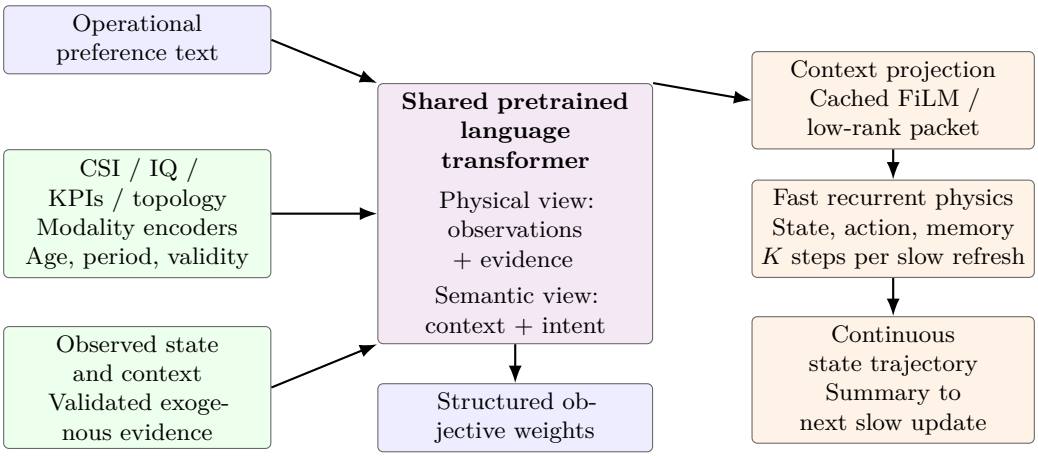


Figure 1: Language-centered architecture implemented by the foundation path. The language transformer is the shared slow backbone; two separate views enforce the preference/evidence boundary. Modality adapters and wireless heads require task training. Fast predictions use the cached packet without calling the language model.

Table 1: Multimodal inputs in the language-foundation integration. Front ends are trainable adapters; their radio representations require task-specific data and alignment.

Input	Representation entering the model	Implemented boundary and limitations
Intent, protocol/evidence text	Language tokenizer IDs and masks	Preferences and exogenous evidence use separate typed fields.
State, traffic, KPIs	Continuous vectors or sequences	Numeric projections preserve values; timestamps and acquisition periods are supplied by the caller.
CSI and IQ	Continuous tokens with real/imaginary feature pairs	Linear/MLP adapters; complex preprocessing and RF calibration are external.
Topology	Node features and optional adjacency	One normalized graph message-passing step before language-backbone fusion.
Visual/map features	Continuous patch feature vectors	Configurable numeric patch adapter; raw-image patch extraction is external.

Qwen2.5-0.5B on CPU, loading 494,032,768 pretrained backbone parameters and connecting CSI, IQ, and topology feature tokens. The five-step output had shape $1 \times 5 \times 6$; numeric-adapter gradients were nonzero with frozen backbone weights, and preference changes left physical predictions unchanged. The record is stored in `pretrained_integration.json`. This establishes software/gradient compatibility rather than a trained RAN model or a general language-understanding score.

The following lightweight variant replaces the expensive language trunk with small train-from-scratch attention/MLP components. Its controlled experiments isolate recurrent dynamics, sharing, and modulation at low cost. The numerical results later in this paper belong *only* to that lightweight variant. They are not evaluations of Qwen or evidence that language pretraining improves wireless prediction.

3 Related Work and Scope

World models learn action-conditioned transitions for prediction and planning; DreamerV3 is an influential example of learning control through a latent world model across diverse environments [7]. MR-NWM is a supervised prediction and finite-candidate auditing study. It does not implement Dreamer, learn an actor, or demonstrate reinforcement-learning performance.

Clockwork recurrent networks demonstrate that recurrent computation can be partitioned by update frequency [2]. HyperNetworks generate parameters for another network [3]; FiLM conditions computation with feature-wise affine transformations [4]. LoRA parameterizes an update using low-rank factors [5]. Here the low-rank factors are generated from each macro context, while the base model is also trained. This differs from the standard frozen-backbone fine-tuning setup of LoRA.

Input-dependent state-space models motivate selective recurrent memory [6]. Our physics expert is a simple gated recurrence, with bounded forgetting coefficients. It is neither an implementation of Mamba nor a physics-constrained channel model. Similarly, the heterogeneous expert mixture executes every expert and therefore offers no sparse-MoE compute saving.

Flow matching regresses a vector field along a prescribed probability path [8]. Transfusion provides an

example of combining discrete and continuous objectives in one multimodal model [9]. We include an optional action-conditioned flow head for future stochastic trajectory studies, but the reported deterministic benchmark does not use it. In particular, no claim about improved uncertainty calibration follows from its presence.

The repository has two explicit realizations. The language-foundation path loads an external pretrained transformer and tokenizer and adds numeric/graph adapters. The lightweight benchmark path uses randomly initialized semantic MLPs, with no pretrained language weights. No over-the-air measurements, RIC integration, or DU deployment are supplied. The controlled-language benchmark cannot support open-domain or foundation-model zero-shot claims.

4 Problem Formulation

4.1 State, actions, context, and intent

Let $s_t \in \mathbb{R}^{d_s}$ denote the observable network state, $a_t \in \mathbb{R}^{d_a}$ a candidate action, and $c_T \in \mathbb{R}^{d_c}$ measured exogenous context for macro window T. Each window contains K micro steps. Operational intent I_T determines objective weights $w(I_T)$ and, in a fuller system, explicit constraints. A causal transition model has the form

$$s_{t+1} = F(s_t, a_t, c_T, \epsilon_t), \qquad U_t = w(I_T)^\top \psi(s_{t+1}, a_t). \tag{5}$$

The preference part of an intent does not directly enter F. If a language input reports an external event or a configuration change that affects dynamics, that information must be extracted into a distinct, validated context variable. If an intent changes the policy, it can of course change the realized future *through actions*. Equation (5) concerns predictions under the same intervened action sequence and context.

This separation gives a testable invariant:

$$\hat{s}_{1:H}(I, a_{0:H-1}, c, s_0) = \hat{s}_{1:H}(I', a_{0:H-1}, c, s_0). \tag{6}$$

The reference simulator satisfies this invariant by construction. A physical event described in text is not interchangeable with an optimization preference; keeping them in a single undifferentiated prompt would make the intended causal semantics ambiguous.

4.2 Time scales and placement

Our notation describes computational rates, not measured radio deadlines. The O-RAN architecture specifies Near-RT RIC control loops on the order of 10 ms to 1 s and permits faster control inside E2 nodes; loop timing is use-case dependent [1]. A sub-millisecond fast model is a possible engineering target for selected functions, rather than a consequence of assigning a module to a DU. Likewise, K is a modeling hyperparameter here; no micro-step in the toy simulator is calibrated to a transmission interval.

5 Lightweight Reference Variant and Training

5.1 Slow heterogeneous mixture

At each macro step, token embeddings for intent, a state projection, a context projection, and the previous fast summary form

$$Z_T = [E_I(I_T); E_s(s_{T,0}); E_c(c_T); H_{T-1}]. \tag{7}$$

The intent tokenizer is fitted only to the controlled training corpus. Learned positions are added to intent tokens; physical token types have their own encoders. A single attention block produces h_i, followed by a dense mixture of one semantic MLP expert E_{sem} and one recurrent expert E_{phy}. For a modality indicator m_i,

$$\pi_i = \text{softmax}(W_g[h_i; \text{onehot}(m_i)]), \tag{8}$$

$$y_i = \pi_{i,\text{sem}} E_{\text{sem}}(h_i) + \pi_{i,\text{phy}} E_{\text{phy}}(h_i). \tag{9}$$

For intent tokens the physics term is the identity, and the recurrent physics expert starts afresh at the first physical token. Physical attention queries are masked from intent keys. Both restrictions are necessary: attention masking alone would still allow a recurrent expert to carry semantic information into physical tokens. The semantic stream may read the physical context.

The mean physical output gives h_T^p. A masked mean over intent outputs gives h_T^I, and a head predicts $\hat{w}_T = \text{softmax}(W_I h_T^I)$ over throughput, queue, and energy features. The router is weakly regularized by

$$\mathcal{L}_{\text{route}} = -\frac{1}{|\mathcal{V}|} \sum_{i \in \mathcal{V}} \log \pi_{i,m_i}, \tag{10}$$

where $\mathcal{V}$ excludes padding. This encourages modality specialization; it does not establish that the experts learn interpretable physical laws.

5.2 Fast recurrence and modulation

A hyper-projection from h_T^p generates $(\gamma_T, \beta_T, A_T, B_T)$, with $\gamma_T = 1 + 0.1 \tanh(\cdot)$, $\beta_T = 0.1 \tanh(\cdot)$, and bounded factors $A_T, B_T \in \mathbb{R}^{d \times r}$. Within a macro window, these tensors are held fixed. Each micro step encodes the current state, candidate action, and measured context:

$$x_k = \gamma_T \odot \text{SiLU}(W_x[s_k; a_k; c_T]) + \beta_T. \tag{11}$$

Let $n_k = \text{LN}(x_k)$ and $(u_k, g_k) = W_{\text{in}} n_k$, with a SiLU applied to u_k. The physics memory is

$$\alpha_k = \sigma(W_\alpha u_k), \tag{12}$$

$$z_k = \alpha_k \odot z_{k-1} + (1 - \alpha_k) \odot \tanh(W_B u_k), \tag{13}$$

$$o_k = x_k + W_o[W_C z_k \odot \sigma(g_k)] + \frac{n_k A_T B_T^\top}{\sqrt{r}}, \tag{14}$$

$$\hat{s}_{k+1} = \hat{s}_k + W_\Delta o_k. \tag{15}$$

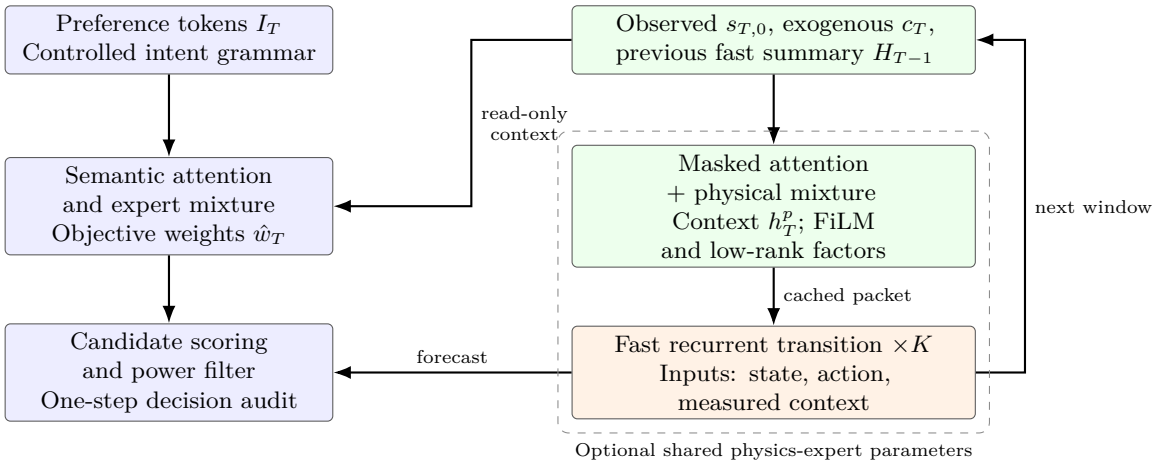


Figure 2: Lightweight reference-variant information flow. Physical prediction has no path from preference tokens. The semantic stream can use physical observations; numerical action evaluation uses the forecast and compiled objectives. Slow and fast execution are separated in the API, while shared training parameters remain a research option.

The last expression is a residual *state increment*; the measured state is not replaced by an action-independent macro forecast. Although z_k has a bounded convex update, this alone does not guarantee stability of the complete autoregressive predictor or the downstream control loop.

In the shared variant, the slow physical expert and fast recurrence refer to the identical parameter object. No extra copy is trained. A separately deployed process would instead receive a versioned copy of the trained parameters; weight sharing during training does not create cross-machine shared memory. The unshared variant allocates an independent fast core. All variants retain explicit dependence on actions and continuous context.

5.3 Feedback and training graph

The K fast hidden states are compressed by a kernel-K, stride-K one-dimensional convolution into H_T. This summary is available to the *next* macro step. It cannot contain future observations. Memory z resets at a macro boundary; information persists through H_T and the terminal predicted state.

Training uses two consecutive macro windows, initialized once with the observed state. At the boundary,

$$s_{T+1,0} = \hat{s}_{T,K}, \qquad H_T = R(o_{T,1:K}). \tag{16}$$

No target state is substituted. Within each macro window, s_k and z_k are detached every $L = 3$ steps. The modulation tensors retain their gradients. The direct path from each micro-step loss to the slow context therefore remains available despite truncation. Shared parameters also receive gradients through the fast recurrence itself. Truncation limits the differentiated history; it neither proves convergence nor prevents all gradient explosions. We additionally clip the global gradient norm to one.

With $H = 2K$, the reference loss is

$$\mathcal{L} = \mathcal{L}_{\text{dyn}} + 0.2\mathcal{L}_I + 0.01\mathcal{L}_{\text{route}}, \tag{17}$$

$$\mathcal{L}_{\text{dyn}} = \frac{1}{BHd_s} \sum_{b,k,j} (\hat{s}_{b,k,j} - s_{b,k,j})^2, \tag{18}$$

$$\mathcal{L}_I = -\frac{1}{2B} \sum_{b,T,j} w_{b,j} \log \hat{w}_{b,T,j}. \tag{19}$$

Semantic targets are determined by the actual words in each training phrase. There is no independently sampled next-token target and no claim of language model pretraining. The optional flow loss is disabled for all reported runs.

For a controlled gradient ablation, the detached variant uses an independent fast core and stopgrad(h_T^p) before its modulator. The modulator, fast core, semantic head, and route loss remain trainable. Comparing this variant with the *unshared* jointly trained model isolates the bridge-gradient choice. Comparing it only with the shared model would also change parameter sharing.

5.4 Optional probabilistic trajectory head

For an optional extension, let $y = s_{T,1:K}$ and $\xi \sim \mathcal{N}(0, I)$. A velocity network conditioned on both h_T^p and the *candidate action sequence* minimizes

$$\mathbb{E}_{\tau,\xi} \left\| v_\phi((1-\tau)\xi + \tau y, \tau, h_T^p, a_{0:K-1}) - (y - \xi) \right\|^2. \tag{20}$$

The sampling API requires an explicit base draw and integrates the velocity using Euler steps. Starting every sample at zero would not reproduce the Gaussian-base generative model. The optional output is separate from Equation (15); it is not silently injected into the deterministic rollout. We test differentiation and sampling reproducibility, while leaving probabilistic accuracy and calibration for a dedicated study.

6 Reproducible Evaluation Protocol

6.1 Normalized two-link environment

The state contains two link gains g, two queues q, and two interference variables j. Actions are two normalized powers $p \in [0,1]^2$; context contains arrival rates λ_1, λ_2, cross-coupling χ, and mobility μ. For each link i, with $\bar{i}$ its neighbor,

$$g_i' = \mathrm{clip}((1-\mu)g_i + 0.55\mu + \epsilon_i, 0.05, 1.2), \quad (21)$$

$$r_i = \log_2\left(1 + \frac{4 g_i p_i}{0.2 + j_i + \chi p_{\bar{i}}}\right), \quad (22)$$

$$q_i' = \mathrm{clip}(q_i + \lambda_i - 0.12 r_i, 0, 2), \quad (23)$$

$$j_i' = 0.65 j_i + 0.35 \chi p_{\bar{i}}, \quad (24)$$

where $\epsilon_i \sim \mathcal{N}(0, 0.005^2)$. Powers are independently sampled uniformly on $[0,1]$. Initial gains, queues, and interference are uniformly sampled from $[0.2, 0.9]$, $[0.1, 0.9]$, and $[0.05, 0.25]$, respectively. Training context uses $\lambda_i \sim U(0.05, 0.15)$, $\chi \sim U(0.08, 0.25)$, and $\mu \sim U(0.03, 0.12)$. The joint shift adds 0.12, 0.30, and 0.25 to those variables, respectively. Thus OOD context values for coupling and mobility do not overlap training.

The system has controllable queues and delayed interference, but no waveform, antenna array, coding, beam management, scheduler, packet simulator, or calibrated propagation model. It is an intentionally small system-identification test. The coefficients and state normalization are synthetic choices; they must not be interpreted as measurements of a mobile network.

6.2 Intent splits and decision audit

Three objective names—throughput, latency, and energy—are used in singleton and two-objective phrases. A target puts equal weight on the named objectives. Three training templates vary the wording. Two held-out templates reorder the same known words. An additional test uses all three objectives together, a combination absent during training. No model or tokenizer is fitted to these test phrases. Because this is a deliberately constrained grammar, its scores measure compositional parsing under that grammar, not open-language zero-shot generalization.

For each test state we also enumerate powers from $\{0.1, 0.4, 0.7\}^2$, retaining only candidates with $p_1 + p_2 \le 0.8$. The utility features are normalized post-transition rate, negative mean queue, and negative mean power. The learned model chooses a candidate using predicted features and inferred weights. We measure the regret against the exact simulator's best *feasible candidate*, evaluated with the true intent weights. A second regret calculation substitutes the true weights into model selection to isolate prediction error from parsing error. The audit is deterministic and one step long; it is not a learned policy or an end-to-end closed-loop control experiment. Zero power violations are an expected property of the candidate filter, not evidence of learned safety.

6.3 Baselines and matched comparisons

The eight executable variants are: shared MR-NWM; unshared MR-NWM; shared MR-NWM without FiLM or low-rank modulation; the unshared detached-bridge control; an action/context-conditioned residual MLP; an approximately parameter-matched MLP; a standalone recurrent model; and persistence. The MLPs have two hidden layers of width 64 and 160, respectively, and predict the same state increments. They have no semantic parser, so semantic scores are not invented for that baseline. The standalone recurrent control uses the same fast recurrence and truncation schedule as MR-NWM, with no slow tier or modulation. Parameters and optimization steps are reported rather than claiming exact compute matching. No pure LLM, pretrained wireless foundation model, or industrial controller is represented by a synthetic row in a table.

All trainable variants use 300 AdamW updates, batch size 48, learning rate 0.002, and weight decay 10^{-4}. MR-NWM uses width 32, recurrent state size 16, rank 2, four attention heads, and $K = 5$. All models receive identical training trajectories, indexed by seeds $100000+u$ for update u. Initialization seeds are 11, 22, and 33. Evaluation uses 256 trajectories per split with seed 900000, never used for training. Models are evaluated after the fixed final update; no test score selects a checkpoint or changes the budget.

6.4 Metrics and timing

We use signal-energy NMSE,

$$10 \log_{10} \frac{\sum_{b,k,j} (\hat{s}_{b,k,j} - s_{b,k,j})^2}{\sum_{b,k,j} s_{b,k,j}^2}, \quad (25)$$

with lower values better. This is distinct from dividing MSE by centered target variance. Per-step MSE and per-state NMSE accompany the aggregate, since a global score can hide errors in small-scale variables. A paired action audit compares the predicted trajectory difference under a and $1-a$ with the noise-free simulator difference. A predictor ignoring actions has zero effect and therefore an action-effect NMSE of 0 dB.

Intent parsing uses total variation distance $\frac{1}{2}\|\hat{w} - w\|_1$ and cosine similarity. Cosine is not converted into an invented "alignment percentage." For latency, the model is in inference mode; measurements use ten warm-up calls followed by 100 timed batch-one calls, with device synchronization around timed regions. We retain raw timing samples and report p50, p95, and p99. The fast measurement uses a cached modulation packet and a supplied recurrent state. It excludes slow encoding, data acquisition, transport, scheduling, model loading, and actuator execution. It therefore cannot establish a radio control deadline or a deterministic worst-case bound.

7 Results and Interpretation

Table 2 reports measurements from the fixed protocol. The shared model reaches -22.24 ± 0.20 dB on held-out IID trajectories, compared with -14.92 ± 0.35 dB for the approximately parameter-matched MLP and -21.19 ± 0.30 dB for the standalone recurrent model. Thus recurrence is a strong alternative to the extra slow tier in this example. The shared model has 28,031 parameters, whereas the standalone recurrent control has 5,398. The latter should remain a serious baseline when a slow semantic/context tier is not otherwise required by the application.

Under the joint traffic, coupling, and mobility shift, the shared predictor's NMSE degrades to -9.58 ± 0.41 dB. The standalone recurrent model reaches -9.43 ± 0.37 dB. This small difference does not establish a robust out-of-distribution benefit from the nested hierarchy. Persistence is also competitive with the smaller MLP under this shift, illustrating that training accuracy alone is an inadequate model-selection argument for deployment.

Sharing changes the IID mean only modestly relative to the unshared model, whose three-seed variability is substantially larger. Removing modulation reduces IID accuracy, but its OOD mean is close to the full model. Most importantly, the correct bridge comparison—unshared joint versus unshared detached—does not establish a consistent advantage for cross-scale gradients across IID and OOD conditions. The experiment therefore supports the feasibility of the training graph and APIs, rather than a general superiority claim for each proposed coupling.

The paired action-effect NMSE for the shared model is -11.75 ± 0.27 dB, below the action-ignorant persistence control's 0 dB. This provides a direct check that the learned predictor uses candidate actions. It does not by itself establish an intervention effect in a real radio system; the comparison is against the explicitly intervenable synthetic simulator.

Table 3 also reveals substantial parsing error despite the small controlled vocabulary. The shared model's mean TV increases from 0.118 for held-out wording to 0.195 for the unseen three-objective combination. An additional exact lexicon compiler achieves zero TV on all of these restricted grammar splits, as recorded in `intent_baseline.json`. That control requires no neural training and makes clear that the present language task does not justify using a pretrained language model. The learned head is a demonstration of differentiable intent compilation whose robustness needs improvement and evaluation on a less artificial corpus.

The shared model's finite-candidate IID regret is 0.0055 ± 0.0022 in normalized utility units. All selected candidates satisfy the configured power budget because infeasible candidates are removed before ranking. Small regret on this short candidate list should not be confused with high long-run policy return or learned constraint satisfaction.

On the measured CPU host (Python 3.12.14, PyTorch 2.5.1 CPU, one PyTorch thread), the shared model's mean per-run fast p50 is approximately 0.081 ms and its mean per-run fast p99 is approximately 0.291 ms. The slow p50 is approximately 0.591 ms. These are measurements of a tiny eager-mode model, with the fast packet already cached. They are neither a comparison with foundation-model inference nor a worst-case deadline guarantee. Raw timing samples and per-initialization quantiles are preserved, so the host-specific observation can be checked without treating it as a deployed baseband result.

8 Deployment Contract and Further Validation

Both Python realizations separate one slow call from K cached-packet fast calls, which allows later export of the fast subnetwork. No transport or asynchronous runtime is implemented. A distributed implementation would need explicit packet validity, state-normalization metadata, model-version checks, monotonic sequence numbers, expiry handling, and a local fallback when context is stale. Parameters shared during training would be copied consistently into the deployed modules. These are requirements for an integration study, not features verified here.

A credible next evaluation should first replace the synthetic environment with calibrated radio simulation and held-out measured traces. Sionna is one available physical-layer research framework [10]. A topology-aware extension must specify graph inputs and compare with graph-based baselines. The supplied pretrained semantic integration must be task-trained and evaluated with its named model, actual checkpoint and license, a disjoint task corpus, and an explicit conversion from free-form instructions to objective and constraint schemas. The controlled-language baseline here should remain to reveal the marginal value of pretraining.

An expanded stochastic study should compare the optional flow head with deterministic regression, ensembles, and simpler likelihood models, including calibration and action-conditioned coverage under distribution shift. Control experiments must then measure policy return, realized constraint violations, action-ranking regret, and recovery after faults over sustained rollouts. Device-specific profiling must include end-to-end deadlines and tail behavior under load, separately from Python kernel timing. The repository provides CPU-tested code and a synchronized accelerator timing path; no NPU or GPU measurement is claimed by this release.

9 Limitations and Conclusion

The benchmark's small state space, short horizon, known context, random power coverage, and templated language make it useful for debugging but inadequate for ranking general-purpose 6G architectures. Three initializa-

Table 2: Measured signal-energy NMSE over ten consecutive forecast steps. Values are mean ± sample standard deviation over three initialization seeds. All rows use the same held-out physical trajectories for a given dynamics split.

Model	Parameters	IID NMSE (dB)	OOD NMSE (dB)
MR-NWM shared	28,031	-22.24 ± 0.20	-9.58 ± 0.41
MR-NWM unshared	32,815	-21.99 ± 1.13	-9.34 ± 0.86
No modulation	28,031	-21.64 ± 0.15	-9.47 ± 0.88
Detached bridge	32,815	-21.63 ± 0.96	-9.56 ± 0.99
Residual MLP	5,382	-14.36 ± 0.07	-6.33 ± 0.23
Matched-size MLP	28,806	-14.92 ± 0.35	-7.15 ± 0.70
Standalone recurrent	5,398	-21.19 ± 0.30	-9.43 ± 0.37
Persistence	0	-3.67 ± 0.00	-6.87 ± 0.00

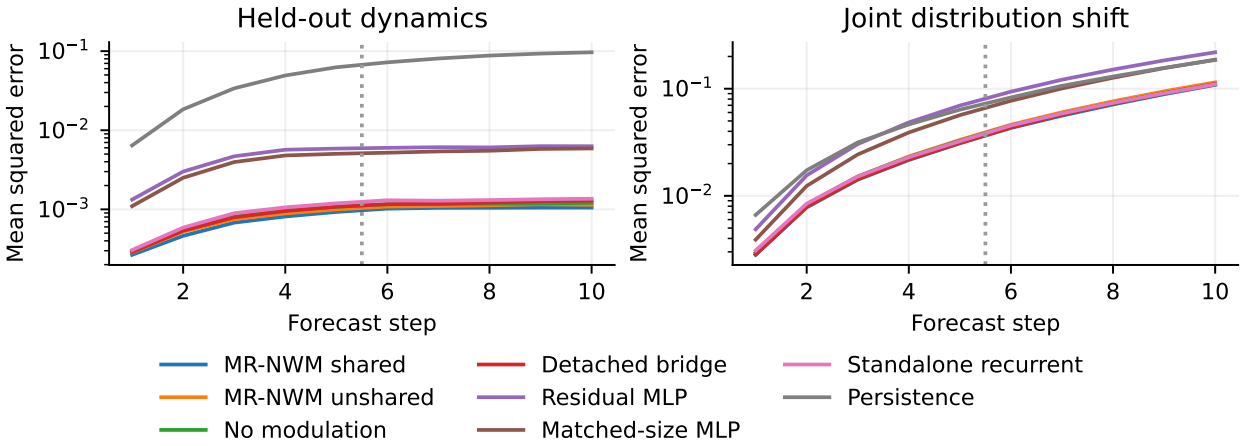


Figure 3: Measured forecast error by horizon. The dotted boundary separates two five-step macro windows; the second window starts from the model's first-window terminal prediction. Curves average three initialization seeds. Raw per-seed curves and per-state NMSE are retained in the result manifest.

Table 3: Controlled-vocabulary parsing and finite-candidate decision audit. Lower is better for all three metrics. Composition refers to the held-out three-objective combination; it is not a foundation-model zero-shot benchmark.

Model	Wording TV	Composition TV	IID regret
MR-NWM shared	0.118 ± 0.040	0.195 ± 0.046	0.0055 ± 0.0022
MR-NWM unshared	0.114 ± 0.040	0.211 ± 0.061	0.0052 ± 0.0028
No modulation	0.114 ± 0.040	0.187 ± 0.053	0.0060 ± 0.0022
Detached bridge	0.111 ± 0.040	0.206 ± 0.056	0.0054 ± 0.0024

tion seeds expose some optimization variation; because evaluation trajectories are shared, their sample standard deviation is not a confidence interval over deployment environments. The synthetic dynamics have no units or calibrated radio timing. The recurrent state update is bounded but the whole predictor is not certified stable. A hard power filter provides only that specific actuator guarantee, not general reliability or safety.

The intended MR-NWM foundation is a pretrained language model with continuous modality adapters and a cached fast physical subnetwork. Within the measured lightweight study, MR-NWM provides an auditable composition of multi-rate prediction, optional physics sharing, modulation, and supervised intent compilation. Its practical value is that the proposed couplings and their limitations are measurable in an executable pipeline. The central empirical question—whether these couplings improve realistic RAN control beyond strong modular baselines—remains open.

Reproducibility and Availability

Source, fixed experiment commands, tests, raw metric and timing samples, source hashes, environment metadata, generated tables, and figure-generation code are available in the repository. The paper tables are generated directly from the committed result JSON. Trained

toy checkpoints are preserved in the release project; their state dictionaries can be re-evaluated using the supplied script. Code is distributed under the included MIT license. The manuscript is a research preprint and makes no claim of peer review.

References

[1] ETSI, "O-RAN Architecture Description," ETSI TS 103 982, V14.0.0, Mar. 2026, Secs. 5.2 and 5.3.2. ETSI specification.

[2] J. Koutník, K. Greff, F. Gomez, and J. Schmidhuber, "A Clockwork RNN," in *ICML*, 2014, arXiv:1402.3511.

[3] D. Ha, A. Dai, and Q. V. Le, "HyperNetworks," in *ICLR*, 2017, arXiv:1609.09106.

[4] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, "FiLM: Visual Reasoning with a General Conditioning Layer," in *AAAI*, 2018, arXiv:1709.07871.

[5] E. J. Hu et al., "LoRA: Low-Rank Adaptation of Large Language Models," in *ICLR*, 2022, arXiv:2106.09685.

[6] A. Gu and T. Dao, "Mamba: Linear-Time Sequence Modeling with Selective State Spaces," arXiv:2312.00752, 2023.

[7] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, "Mastering Diverse Domains through World Models," arXiv:2301.04104, 2023.

[8] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le, "Flow Matching for Generative Modeling," in *ICLR*, 2023, arXiv:2210.02747.

[9] C. Zhou et al., "Transfusion: Predict the Next Token and Diffuse Images with One Multi-Modal Model," arXiv:2408.11039, 2024.

[10] J. Hoydis et al., "Sionna: An Open-Source Library for Next-Generation Physical Layer Research," arXiv:2203.11854, 2022.

[11] Qwen Team, "Qwen2.5-0.5B, pretrained base language model, 2024. `https://huggingface.co/Qwen/Qwen2.5-0.5B`.